



Lecture 4: Priority Queues, Heaps and Heapsort



Prof. Aleksander Mądry

Reading: CLRS 6.1-6.4

Priority Queue (PQ)

An *abstract data structure* (aka *data type*) maintaining a set S of **elements**, each associated with a **key**, supporting the following operations:

$\text{insert}(S, x)$: insert element x into set S

$\text{max}(S)$: return element of S with largest key

$\text{extract_max}(S)$: return element of S with largest key and remove it from S

$\text{increase_key}(S, x, k)$: increase value of x 's key to new value k
(assumed to be \geq the current key value)

Think: All the operations you would need to organize triage in an emergency room \rightarrow key value = severity of patient's condition

(Tons of applications in algorithms and across the whole CS too.)

Priority Queue (PQ)

An *abstract data structure* (aka *data type*) maintaining a set S of **elements**, each associated with a **key**, supporting the following operations:

$\text{insert}(S, x)$: insert element x into set S

$\text{max}(S)$: return element of S with largest key

$\text{extract_max}(S)$: return element of S with largest key and remove it from S

$\text{increase_key}(S, x, k)$: increase value of x 's key to new value k
(assumed to be \geq the current key value)

This abstraction specifies desired functionality/interface,
but how to implement it?

Naïve way: Use an (unsorted) array and scan all elements to find max



→ each insert and increase key takes $\Theta(1)$ time

→ all the other operations take $\Theta(n)$ (worst-case) time

Can we
do better?

(Max) Heap

- **Data structure** implementing a priority queue
- It is an **array** that:
 - ➔ we visualize as a (nearly complete) **binary tree**
 - ➔ satisfies **Max Heap Property (MHP)**:
Key of a node is \geq than the keys of its children
- (**Min Heap** defined analogously)




Important fact:
Height of the tree is
always $O(\log n)$

Mapping Tree to a Heap

- root: first element in the array (**$i=1$**)
- parent(i): **$\text{floor}(i/2)$** returns index of node's parent
- left(i): **$2i$** returns index of node's *left* child
- right(i): **$2i+1$** returns index of node's *right* child

(**Note:** No pointers needed!)

Important detail: We index elements starting from **$i=1$** here




Why Heaps?


Key consequence of Max Heap Property:

Root/first element is always the max → can do $\max(S)$ in $\Theta(1)$ time!
(Note: the array is not sorted though!)

But: How to maintain the Max Heap Property property
after insert/extract_max/increase_key?



In fact:
How to build a Max Heap
to begin with?




Key Primitive

`max_heapify(A[i]):` Corrects a **single** violation of Max Heap Property in a subtree rooted at **i only**

How to implement it?

➔ Assume that the trees rooted at `left(i)` and `right(i)` are Max Heaps

➔ If element `A[i]` violates the MHP, correct violation by “trickling” this element down the tree, making the subtree rooted at `i` a Max Heap (**Important:** Always swap with the larger of two children. **Why?**)



Want this to be a Max Heap


Key Primitive

`max_heapify(A[i]):` Corrects a **single** violation of Max Heap Property in a subtree rooted at **i only**

In other words:

➔ Assume that the trees rooted at `left(i)` and `right(i)` are Max Heaps

➔ If element `A[i]` violates the MHP, correct violation by “trickling” this element down the tree, making the subtree rooted at `i` a Max Heap (**Important:** Always swap with the larger of two children)




Key Primitive

`max_heapify(A[i]):` Corrects a **single** violation of Max Heap Property in a subtree rooted at **i only**

In other words:

➔ Assume that the trees rooted at $\text{left}(i)$ and $\text{right}(i)$ are Max Heaps

➔ If element $A[i]$ violates the MHP, correct violation by “trickling” this element down the tree, making the subtree rooted at i a Max Heap (**Important:** Always swap with the larger of two children)




Run time?

➔ $\Theta(1)$ at each level

➔ total: $O(\text{subtree height})$
 $= \Theta(\log n)$ (worst case)
(Recall: tree is balanced)


Guaranteed to be a Max Heap now

Implementing Extract_Max




Implementing Extract_Max

- Swap the root with the last element of the heap




Implementing Extract_Max




- Swap the root with the last element of the heap
- Now we can remove it from the heap (decrease the heap size by one)
- **To fix MHP:**
Max_heapify the root

Implementing Extract_Max




- Swap the root with the last element of the heap
- Now we can remove it from the heap (decrease the heap size by one)
- **To fix MHP:**
Max_heapify the root

Implementing Extract_Max




- Swap the root with the last element of the heap
- Now we can remove it from the heap (decrease the heap size by one)
- **To fix MHP:**
Max_heapify the root

Implementing Extract_Max



- Swap the root with the last element of the heap
- Now we can remove it from the heap (decrease the heap size by one)
- **To fix MHP:**
Max_heapify the root

Implementing Extract_Max




- Swap the root with the last element of the heap
- Now we can remove it from the heap (decrease the heap size by one)
- **To fix MHP:**
Max_heapify the root
- Done!

Run time?

- ➔ $\Theta(1)$ (swapping) + $\Theta(1)$ (removal) + $O(\log n)$ (max_heapify)
- ➔ total: $\Theta(\log n)$ (worst case)


Implementing Insert

(in a sense: “reversing” the extract_max)



Implementing Insert


(in a sense: “reversing” the extract_max)



- Add the new element as the last one
- **To fix MHP:**
“Promote” the new element up the tree
 (“reversed” max_heapify)

Implementing Insert


(in a sense: “reversing” the extract_max)



- Add the new element as the last one
- **To fix MHP:**
“Promote” the new element up the tree
 (“reversed” max_heapify)

Implementing Insert

(in a sense: “reversing” the extract_max)




- Add the new element as the last one
- **To fix MHP:**
“Promote” the new element up the tree
 (“reversed” max_heapify)
- Done!

Run time?

- ➔ $\Theta(1)$ (addition) + $O(\log n)$ (promotion up the tree)
- ➔ total: $\Theta(\log n)$ (worst case)

Implementing Increase_key (Similar to Insert)




Implementing Increase_key

(Similar to Insert)


- Increase the key value

- **To fix MHP:**

Again, “promote” the new element up the tree




Implementing Increase_key (Similar to Insert)



- Increase the key value
- **To fix MHP:**
Again, “promote” the new element up the tree

Implementing Increase_key (Similar to Insert)



- Increase the key value
- **To fix MHP:**
Again, “promote” the
new element up the tree
Done!

Run time?

- ➔ $\Theta(1)$ (key value increase) + $O(\log n)$ (promotion up the tree)
- ➔ total: $\Theta(\log n)$ (worst case)

How to build a heap from a scratch?

Simple way:

- Start with an empty heap
- Insert all the n elements into it
- Total time: $\Theta(n \log n)$ (worst-case)

Better way:

- Use divide & conquer! (see blackboard)
- $T(n) = 2 T(n/2)$ (conquer) + $O(\log n)$ (in-place divide & combine)
- Total time: $\Theta(n)$ (by Master Theorem)

Iterative (and in-place) way:

build_max_heap(A):


 for $i=n$ downto 1

 do max_heapify(A,i)

$O(n \log n)$

→ Total time? At first glance: ~~$\Theta(n \log n)$~~ (cost of n max_heapify)

Actually: $\Theta(n)$ (see blackboard)



Cool application: Sorting

Heapsort: Sorting using a heap/priority queue

- ➔ Build a heap out of all elements
- ➔ Extract_max all elements one-by-one in an (inversely) sorted order!
- ➔ Total time: $\Theta(n \log n)$ (worst-case)

This is a different algorithm than Merge sort!

In particular: Heapsort is actually an in-place algorithm (once we unravel the implementation of the heap)